A Tensor-Based Expansion of the Cube Rule for Generative Nutrient Morphology

Dr. Jane Lydia Adams, Ph.D.

Abstract

We extend the Cube Rule of Food by using its morphology categories as one axis of a third-order tensor $T \in \{0,1\}^{I \times J \times K}$ whose axes encode cube morphology, protein type, and starch type. Within this ontology, canonical foods can be placed reproducibly, collisions become explicit coordinates with multiple occupants, and empty cells can be interpreted as constrained sites for culinary speculation. We show in particular that salad occupies exactly the `starch_type=none` slice, examine mono-starch edge cases such as plain rice and bowls of croutons as evidence that the Cube Rule may require a dedicated tenth mono-starch category, and characterize sparsity in the tensor as a mix of structural impossibility and uneaten potential. To populate and analyze the space, we combine ontology-constrained LLM generation, deterministic post-generation filtering, occupancy visualizations, and a supporting software pipeline that renders both the paper figures and an interactive local web viewer for the tensor outputs.

1 Introduction

This project studies a formal extension of the Cube Rule of Food, a morphology-based framework originating in Phosphatide's "Grand Unified Theory of Food Identification" [4] and later popularized and expanded by André Arko at `https://cuberule.com/` [3], using a rank-3 tensor T over three categorical axes. The immediate objective is to define a compact ontology in which each food is assigned a coordinate triple (i, j, k) in the discrete latent space $\{1, \ldots, I\} \times \{1, \ldots, J\} \times \{1, \ldots, K\}$, so that canonical Cube Rule examples are representable and ambiguous cases are explicit. Our framing treats occupancy, sparsity, and ambiguity as geometric properties of T rather than anecdotal judgments. Under this formulation, cuisine is represented as a discrete structure over the index space rather than as an informal list of examples.

1.1 The Cube Rule of Food

The Cube Rule of Food is a morphology-based classification scheme that assigns a dish according to where structural starch appears on the six faces of an idealized cube. The core idea originates with Phosphatide's "Grand Unified Theory of Food Identification" [4], while the Cube Rule formulation used here follows the later presentation and elaboration by André Arko [3]. In the original Rule, these labels are morphology-only categories rather than cuisine names: *salad* has no structural starch faces, *toast* has a bottom starch face, *sandwich* has top and bottom starch faces, *taco* has left, right, and bottom starch faces, *sushi* has top, left, right, and bottom starch faces, *quiche* has side walls and a bottom with an open top, *calzone* is fully enclosed, *cake* consists of stacked starch layers, and *nachos* denotes an outer container plus interior starch pieces. Some of the Rule's most memorable examples are notably surprising: nigiri counts as *toast*, an Uncrustable counts as *calzone*, and a salad with croutons counts as *nachos*; for the lettuce-crouton boundary case discussed later, see Figure 3. These nine morphologies define the cube-type axis used throughout this paper and supply the structural vocabulary that the tensor formalism extends with protein and starch-type axes. Figure 1 adapts the Cube Rule morphology

diagram in schematic form from Arko's presentation [3], itself an adaptation and popularization of Phosphatide's original post [4].

2 Related Work

This work builds on the Cube Rule as a prior food-classification framework and introduces two external strands of literature to formalize the proposed tensor-based approach. The first strand is modern LLM control and prompting: prompt taxonomies, schema-constrained generation, grammar-constrained decoding, guided decoding, and retrieval-augmented generation provide language for the prompt design, structured decoding, exemplar use, and post-generation filtering described in the Methods section. In this paper those ideas appear concretely when the model is asked for a dish matching a specific cell such as (toast, seafood, rice) and is steered toward outputs like nigiri rather than loosely related rice dishes or free-form descriptions. The second strand is tensor and geometry-oriented representation theory: tensor completion, manifold hypotheses, simplicial complex constructions, persistent homology, and topological data analysis provide language for discussing empty-cell completion, clustered occupancy structure, and topological features of the support set $\mathcal{S} = \{(i,j,k) : T_{i,j,k} = 1\}$ in higher-order categorical space. Here the motivating examples are food-level rather than abstract: filling an empty cell with a plausible named dish is treated as tensor completion, while repeated occupancies around nearby dishes such as nigiri, sushi rolls, and ramen motivate the geometric language for local structure. Classical combinatorics references are also relevant for occupancy guarantees, including direct pigeonhole-style collision arguments when the number of foods exceeds $|I \times J \times K| = IJK$. Taken together, these literatures provide the methodological and conceptual vocabulary used to define, populate, and analyze the proposed taxonomy.

3 Methods

We define two state spaces in direct analogy to the real numbers $\mathbb{R}$ (the space of all real numbers). Let $\mathbb{F}_{\text{real}}$ denote the *theoretical* set of all *real* foods: those taken to have been cooked and recorded in history (the culinary analogue of "real" in $\mathbb{R}$); we approximate this set with our catalog rather than observe it in full. Let $\mathbb{F}_\infty$ denote the *theoretical* set of all *conceivable* foods, including in principle dishes that have never been prepared—the full state space of possible foods, which we cannot directly verify—with $\mathbb{F}_{\text{real}} \subseteq \mathbb{F}_\infty$. Our tensor formalism maps elements of $\mathbb{F}_\infty$ to coordinates; the catalog we maintain is a finite subset of $\mathbb{F}_{\text{real}}$, and we speculate over elements of $\mathbb{F}_\infty \setminus \mathbb{F}_{\text{real}}$ for empty cells.

Let $T \in \{0,1\}^{I \times J \times K}$ denote a tensor indexed by cube type (axis $i \in \{1,\ldots,I\}$), protein class (axis $j \in \{1,\ldots,J\}$), and starch type (axis $k \in \{1,\ldots,K\}$). The starch axis includes a distinguished `none` value for dishes with no structural starch substrate. We set $T_{i,j,k} = 1$ when at least one food maps to index triple (i,j,k); the *support set* (or, in less restrained prose, the Standard Meal Set) is $\mathcal{S} = \{(i,j,k) : T_{i,j,k} = 1\}$ and the set of *empty cells* is $\mathcal{E} = \{(i,j,k) : T_{i,j,k} = 0\}$. In implementation, each occupied cell may contain multiple candidate foods, so the data pipeline stores per-food records and derives occupancy from those records; formally we retain binary $T_{i,j,k}$ while allowing multiple foods per (i,j,k) in the database. This can be interpreted as a categorical incidence tensor with sparse $|\mathcal{S}| \ll IJK$: for example, cells for nigiri, burritos, ramen, and cheesecake may be occupied while many nearby combinations remain empty. Operationally, this turns informal culinary intuition into an explicit indexed representation.

Under the Cube Rule, salad morphology means no structural starch faces. In the tensor, such cases occupy the slice $(i_{\text{salad}}, j, k_{\text{none}})$. A dish that contains structural starch (e.g., a salad with croutons) is reclassified as nachos (outer container plus inner starch), not salad. Admissible cells therefore satisfy the compatibility condition

$$i = i_{\text{salad}} \iff k = k_{\text{none}}.$$

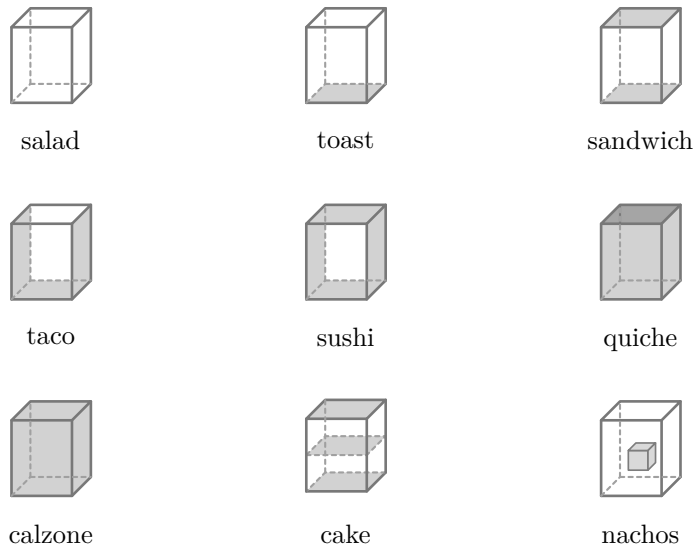

Figure 1: Cube-type morphology schematics for the nine categories on axis i, adapted from Arko's Cube Rule presentation [3], itself adapted from Phosphatide's original "Grand Unified Theory of Food Identification" post [4]. The shaded regions indicate structural starch placement; the drawings are schematic rather than literal geometric reconstructions.

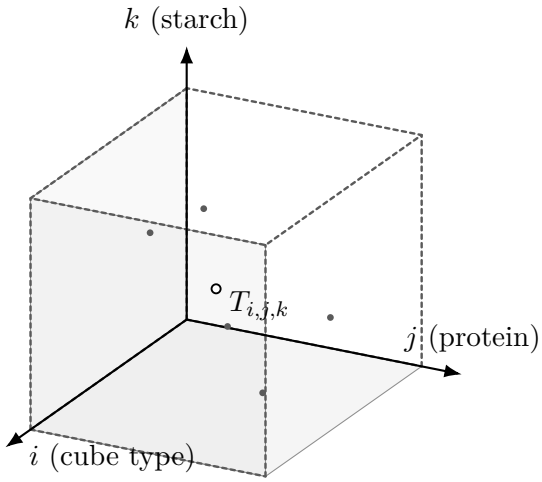

Figure 2: Tensor schematic. The cube-type, protein, and starch axes define the index space of T; filled markers illustrate occupied cells in $\mathcal{S}$, and the highlighted marker indicates a generic cell (i, j, k).

In words, a cell has salad morphology if and only if it has `starch_type=none`. Figure 1 in the Introduction summarizes the nine cube-type morphologies used on axis i.

3.1 The Rice Exception, Mono-Foods, and Dish-Level Classification

The Cube Rule explicitly declines to rule on one case: "You are free to interpret the nature of rice however you wish." It is the only food for which the Rule gives no cube ruling. The Rule *does* address rice in context: nigiri is toast, sushi rolls are sushi, ramen is nachos. The ambiguity is thus not rice-in-dishes but *plain rice by itself*—rice as a starch-based mono-food. The Rule's nachos examples (couscous, poutine, ramen) are illustrated with other things mixed in (couscous with add-ins, poutine with cheese and gravy); nachos is thus starch *plus* other components, not a bowl of plain starch alone. We therefore speculate that the exception applies to the nature of *plain rice alone*, and that a similar ambiguity holds for other starch-based mono-foods: a bowl of plain pita chips, a bowl of bread crumbs—single-starch items with no second component.

Definition 1 (Measure-zero carbs)**.** *A mono-starch food occupies a set of culinary measure zero in the space of toppings, but remains topologically nontrivial.*

If one follows the Rule's existing commitments literally, then plain starches by themselves appear to behave like *salads*: mashed potatoes are classified as salad (with soup treated as wet salad), and a similar logic would make a bowl of croutons salad as well. This yields a gradient from salad to nachos with the addition of more things: lettuce

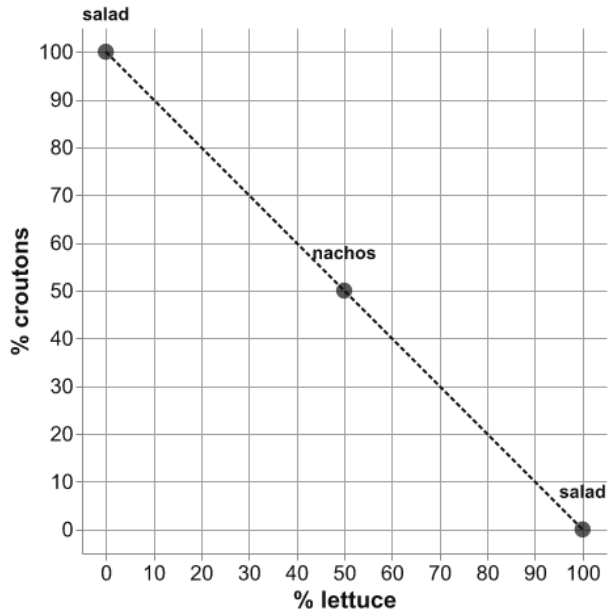

Figure 3: Horseshoe theory of salad: percent lettuce vs. percent croutons. The extremes $(0, 100)$ and $(100, 0)$ are salad (single component); the middle $(50, 50)$ is nachos. Dotted line: lettuce-crouton proportion gradient.

with no croutons is salad; lettuce with croutons is nachos; but a bowl of croutons alone reverts to salad because it is again a single food. Here and throughout, when we discuss mono-food ambiguity we mean starch-based mono-foods such as plain rice or a bowl of croutons; starchless mono-foods are already salad under the ontology (with soup treated as wet salad). In analogy to the horseshoe theory of political economy—where the far left and far right are argued to resemble each other more than either resembles the center [1]—we suggest this be considered the horseshoe theory of salad: the extremes (all lettuce or all croutons) are both salad, while the middle (lettuce and croutons together) is nachos (Figures 3 and 4); in the standard lettuce-centered telling, this is simply the Horseshoe Theory of Lettuce. For the purposes of this paper, we retain the Rule's current dish-level categories and use the horseshoe construction only to describe how mono-starch edge cases behave under the existing ontology. We therefore leave that repair on the back burner and return in the Discussion to whether these cases motivate a new category beyond the present nine.

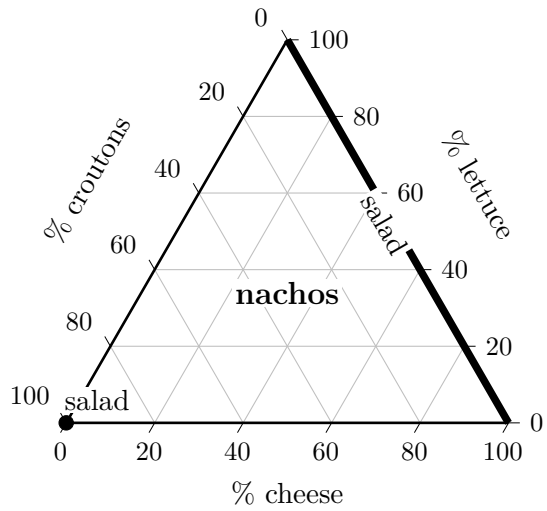

Figure 4: Ternary extension of the horseshoe theory of salad over lettuce, croutons, and cheese. The highlighted crouton-free, lettuce-cheese edge is the ternary trace of the keto plane: any combination of lettuce and cheese with croutons fixed at zero remains salad rather than nachos (assuming no new starch is added, such as quinoa). Note that the whole interior region is considered nachos, because any amount of croutons morphologically introduces a small amount of starch enclosed by the outer container. However, in keeping with the logic that mashed potatoes are salad, we consider a dish that is 100% comprised of croutons to also be salad.

3.2 Constrained LLM Candidate Generation

Cells $(i, j, k) \in \mathcal{E}$ (candidates in $\mathbb{F}_\infty \setminus \mathbb{F}_{\text{real}}$) are queried with a morphology-first prompt that includes a compact Cube Rule crash course, type-specific definitions, and canonical Cube Rule examples. This prompt can be viewed as a small prompt taxonomy: it separates structural morphology instructions, axis-definition instructions, and admissibility constraints so that each part of the classification problem is stated explicitly. The prompt enforces exact axis matching for cube morphology, protein class, and starch class rather than semantic similarity. For example, bottom-face-only starch placements are treated as toast morphology even when the dish name is not colloquially "toast". The API call uses schema-constrained generation rather than free-form text parsing: candi-

dates are returned through a tool payload with typed fields (`shortname`, `description`, `is_real`, `confidence`, and `rationale_brief`). Operationally, this is also a practical form of grammar-constrained decoding, since only outputs that satisfy the required field structure can be returned; a candidate such as "cheesecake" must arrive as a typed row for a specific cell rather than as an essay about dessert. Prompt instructions also prioritize commonly recognized dish names over compositional labels; when only weak descriptive names are available, the model is instructed to return no candidate. The inclusion of canonical Cube Rule examples functions as lightweight retrieval-augmented generation: the model is not asked to classify from an unconstrained prior, but from an in-prompt bank of morphology-defining exemplars such as pizza slice, nigiri sushi, burrito, and lasagna. Prompt scaffolding also states that cube-type labels are morphology-only tags: for example, "salad" does not imply a cold dish, and "taco" does not imply spice or cuisine identity. Protein semantics are also made explicit in prompting and filtering: `none` is reserved for cases without meaningful protein contribution, while dairy ingredients (for example milk, cream, yogurt, and cheese) are treated as dairy-protein signals rather than `none`. Accordingly, the canonical Cube Rule flapjacks example is placed at $(i, j, k) = (\texttt{cake}, \texttt{dairy}, \texttt{wheat})$ rather than $(\texttt{toast}, \texttt{none}, \texttt{wheat})$, because the butter sandwiched between layers contributes a meaningful dairy signal; by contrast, a toast sandwich is still a canonical *sandwich* example under the Rule rather than *cake*, since its structure has top and bottom starch faces instead of stacked starch layers with other materials in between [6].

3.3 Model and Cost Disclosure

All LLM-assisted generation runs used Anthropic Claude (model: `claude-opus-4-6`). The total API cost for the full set of reported experiments was $5.77, generously supported by a grant from the author.

3.4 Hard Validity Constraints

The generation prompt includes hard constraints for strict matching:

- hard axis-locking on starch type: no starch substitution across classes, so a `pastry_dough` cell may admit key lime pie or a Pop-Tart but not falafel pita, quesadilla, or other wheat/corn near-misses;
- hard axis-locking on protein type: no protein substitution across classes, so a `red_meat` cell may admit spaghetti with meatballs or a burrito but not a falafel wrap, cheesecake, or egg custard tart;
- single-entity decoding constraint: no composed "plate" outputs that combine separate foods (for example, "steak with mashed potatoes" or "soup served with bread"); the candidate must be one integrated dish rather than a menu pairing;
- conservative abstention behavior: if no obvious match exists, the model should return an empty candidate list rather than rationalize a weak fit or invent a strained dish name for the sake of filling the cell.

Additionally, cube/starch pairings that violate the structural ontology are prohibited: salad is generated only in the `starch_type=none` slice, and non-salad morphologies are generated only with non-`none` starch types.

These same constraints are enforced after generation using deterministic lexical guardrails, yielding a lightweight constraint-satisfaction layer over candidate rows. In that sense, the pipeline combines stochastic proposal with guided decoding: the model is steered both at generation time and at acceptance time toward outputs that satisfy the ontology, so structurally invalid salad+starch proposals and plate-style outputs are vetoed before review. Candidates failing these checks are rejected before they enter the LLM candidate table for downstream review. Additionally, candidates below a configured confidence threshold are dropped so uncertain guesses do not set $T_{i,j,k} \leftarrow 1$

for speculative (i, j, k). In brief, candidate proposal is stochastic, but acceptance is governed by deterministic constraints and review.

3.5 Ground-Truth Calibration

We validate the filtering pipeline against a ground-truth set derived from canonical Cube Rule examples. The validation script checks axis mapping consistency and verifies that canonical Cube Rule examples pass the same lexical and protein/starch guardrails used for LLM-generated rows. During calibration, several false negatives were identified and corrected by refining token matching and conflict rules (for example, avoiding substring-only matches and reducing over-restrictive composition heuristics). After these adjustments, all canonical Cube Rule examples pass validation under the active guardrail configuration.

3.6 Visualization Axis Ordering

For occupancy visualizations, the order of indices on each axis (e.g., which i appears left-to-right) is computed with hierarchical clustering rather than fixed ontology order. Each index is represented by its binary occupancy profile (e.g., for fixed i, the vector $(T_{i,1,1}, \ldots, T_{i,J,K})$ flattened); pairwise dissimilarity is computed with Jaccard distance, and leaves are ordered by average-linkage dendrogram seriation. This ordering groups indices with similar occupancy structure and is used consistently in both pairwise projections (e.g., $T_{i,j,\cdot}$ over i, j) and 3D categorical axes, producing an interpretable occupancy-manifold view of the tensor slices. At a descriptive level, this reflects a weak manifold hypothesis: foods that are nearby in ontology space, such as nigiri and sushi rolls or salad and nachos edge cases involving croutons, tend to exhibit related occupancy patterns rather than arbitrary dispersion across the tensor.

4 Software Contribution

This project includes a full software pipeline for the paper: scripts to ingest and validate canonical Cube Rule examples, generate candidate foods for empty tensor cells with an LLM, filter and review those candidates, render static and interactive tensor figures, and compile the manuscript PDF. It also includes a small local web application that serves the generated interactive figures through a browser interface, so the tensor projections, 3D occupancy views, and supporting figures can be inspected outside the paper PDF. In that sense, the codebase is not merely an appendix to the paper; it is the executable form of the ontology, the data-processing pipeline, and the figure-viewing workflow described in the surrounding sections.

5 Experiments

The experiment workflow has three steps: map observed foods to coordinates (i, j, k) and set $T_{i,j,k} = 1$ accordingly; measure occupancy $|\mathcal{S}|$ and collisions (multiple foods at the same (i, j, k)); and generate candidate foods for cells in $\mathcal{E}$. The third step can be interpreted as a constrained categorical form of tensor completion: we do not impute arbitrary values, but instead propose named dishes for previously unoccupied cells under hard ontological constraints, such as asking for a plausible dish at a previously empty cube/protein/starch combination rather than assigning an anonymous latent value. Initial analysis reports which regions of the index set $\{1, \ldots, I\} \times \{1, \ldots, J\} \times \{1, \ldots, K\}$ are dense ($T_{i,j,k} = 1$ for many (i, j, k) in a neighborhood), sparse, or structurally inconsistent under the chosen ontology; the sparse regions can also be read as zones of uneaten potential, places where the ontology's menu is still thin. From a representation-learning perspective, this is a small-scale study of combinatorial generalization under hard axis constraints on i, j, and k: the model must recombine familiar ingredients and morphologies into specific candidates such as a new toast-type seafood/rice dish or a new quiche-type dairy/pastry dish. Empirically, the workflow tests whether culinary naming conventions can be mapped consistently onto the proposed coordinate system. Axis categories in Figure 5 are ordered by occupancy-profile hierarchical clustering as described in the Methods section (applied to the indices i, j, k along each axis).

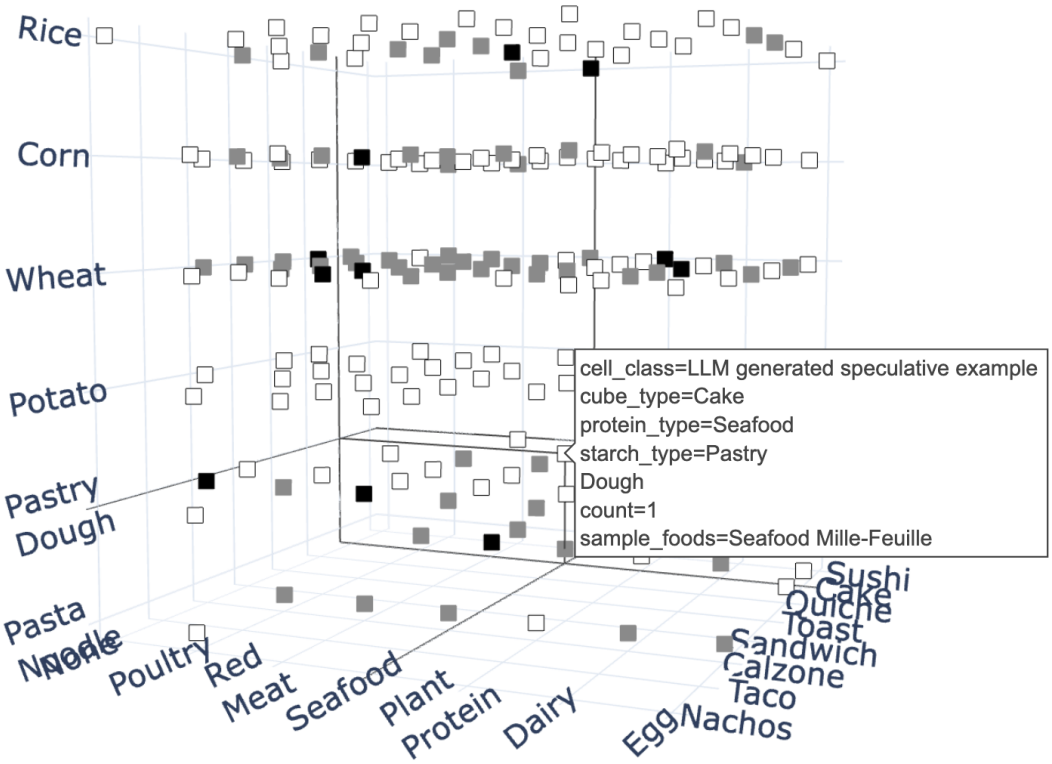

Figure 5: Screenshot of the interactive 3D occupancy tensor. Here we see $T(\text{cake}, \text{seafood}, \text{pastry_dough})$, a speculative LLM-generated example entitled "Seafood Mille-Feuille." This is a savory take on a traditionally sweet French dessert filled with crème pâtissière, and a smoked-salmon mille-feuille variant indeed exists [7].

Review outcomes. All LLM-generated candidates exported for review were accepted and included in the real-only tensor views; the distribution of model confidence for those accepted candidates is shown in Figure 6. These accepted rows can therefore be read as reviewed tensor-completion proposals rather than unconstrained generative samples.

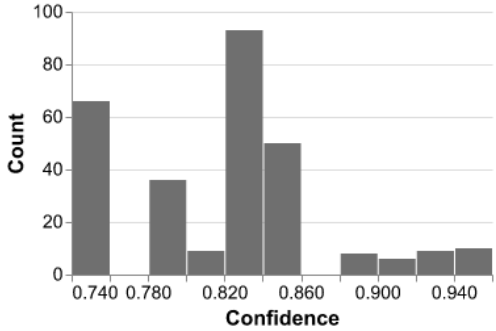

Figure 6: Distribution of model confidence for accepted LLM-generated candidates. The shape is visibly non-parametric because these confidence values are bounded self-reports on $[0, 1]$, then truncated by the acceptance cutoff at 0.75, so the observed mass piles up at a few high-confidence values rather than following a Gaussian law.

5.1 The Candy Salad Conundrum

We also run a targeted analysis of foods whose names explicitly contain "salad" (for example, chicken salad, egg salad, Snickers salad[1], and jello salad[2]). For each item, we ask the model

[1]Snickers salad is a dessert-style potluck dish, commonly made with chopped Snickers bars, tart apples, and whipped topping (often with pudding), and is widely associated with Upper Midwest potluck and church-supper traditions (especially Minnesota and Iowa), though no single point of origin is well documented.

[2]Jello salad is a molded gelatin dish, often containing fruit, marshmallows, or whipped topping, that appears in church-supper, holiday, and potluck traditions across the Upper Midwest and other parts of the United States; despite the name, its structure is far removed from the vegetable-forward sense of "salad."

to classify structural starch placement and thus assign a cube index i (and implicitly j, k). Table 1 reports all tested items, including both foods that satisfy the Cube Rule salad definition and foods that violate it; in the tensor, the former correspond to the 'salad' morphology with `starch_type=none`. We test the hypothesis that many such violations arise from discrete starch inclusions (for example, croutons, pasta pieces, or candy chunks), which shift morphology from the salad index i_{salad} toward the Cube Rule "nachos" index i_{nachos} (outer dish plus inner starch).

6 Discussion

6.1 Collisions and Ontology Size

This is mathematically trivial but operationally useful: once the catalog grows past the available cell count, collisions cease to be a modeling accident and become a structural consequence of the ontology. At that point the ontology begins to boil over: the problem is no longer too many cooks in the kitchen, but too many dishes on the same plate. This motivates optional extensions with additional axes (e.g., a fourth index ℓ), which can be interpreted as moving from $T \in \{0,1\}^{I \times J \times K}$ toward a higher-order multiway representation. For example, future work could introduce additional dimensions such as vegetable type (categorical) or spice level (ordinal or quantitative, depending on how pungency is formalized), but doing so would require explicit decisions about data type and measurement scale before such axes could be used to disambiguate colliding items. One especially promising extension is a cuisine-type or cultural-origin axis. Motivated in part by the informal cross-cultural observation we will call the "Every Country Has a Dumpling Theory" [5], such an axis could help identify morphologically similar (i, j, k) combinations that are common in some culinary traditions but absent in others. It could also surface cross-cultural analogies and speculative transfers more carefully, provided that the model does not treat religious dietary constraints or other culturally meaningful prohibitions as ordinary missing data to be imputed away. One practical tension is model granularity: under-specified ontologies (small I, J, K) collapse semantically distinct foods into the same (i, j, k), while over-specified ontologies create extreme sparsity ($|\mathcal{S}|/IJK$ very small) and weak empirical support per coordinate. The challenge is to add enough garnish to separate meaningful distinctions without overseasoning the ontology into extreme sparsity. That same trade-off appears inside existing axes: the current `other_grain` starch bucket is convenient, but expanding it to distinguish grains such as teff and quinoa could reduce ambiguity and reveal new neighborhoods of the tensor. Such a refinement might separate established regional foods more cleanly while also making room for plausible cross-category completions, for example an injera burrito built around a teff-based wrapper rather than treating all non-wheat, non-rice grains as the same starch substrate.

6.2 Scope of the Ontology

Future extensions should resolve boundary cases where sugar-dominant processed inclusions (for example nougat/candy in snickers salad) may trigger non-salad intuitions without mapping cleanly to the current starch index k. This paper restricts attention to edible foods, even though the original Cube Rule also discusses non-food cases and famously proposes, for example, that humans are ravioli (that is, calzones under the Rule's morphology). An expanded ontology with additional starch and protein categories might therefore reveal interesting findings outside ordinary cuisine as well: a leatherbound book, for instance, could be modeled as an inverse taco in which a protein-like exterior surrounds a filling of starch. We also restrict the morphological axis to ordinary three-dimensional food geometry and do not attempt higher-dimensional morphology. Future work could ask whether tesseract-like foods or other higher-order culinary solids require an additional extension of the cube axis itself, effectively promoting morphology to a richer higher-dimensional index.

Food	Salad?	Cube Type	Rationale
chicken salad	yes	salad	No meaningful structural starch component; bound protein mixture without starch faces or interior starch pieces.
egg salad	yes	salad	No structural starch component; egg-and-dressing mixture does not place starch on any cube face.
tuna salad	yes	salad	Protein mixture with dressing but no starch substrate, shell, or internal starch pieces.
ham salad	yes	salad	Chopped ham in a dressed mixture with no meaningful structural starch component.
potato salad	no	nachos	Discrete starch pieces (potato) mixed into a dressed base. Inner-starch pattern → nachos.
macaroni salad	no	nachos	Discrete pasta pieces mixed with dressing/veg. Inner-starch pattern → nachos.
pasta salad	no	nachos	Discrete pasta pieces mixed with vegetables and dressing. Inner-starch pattern → nachos.
fruit salad	yes	salad	Mixed fruit with no structural starch substrate, enclosure, or layered starch component.
broccoli salad	yes	salad	Broccoli-based dressed mixture; typical add-ins do not create a structural starch pattern.
jello salad	yes	salad	Molded gelatin may be unusual, but it does not introduce structural starch faces under the Cube Rule ontology.
ambrosia salad	yes	salad	Fruit-marshmallow-whipped mixture lacks a distinct structural starch component.
snickers salad	no*	nachos	Discrete candy chunks mixed into a creamy base. Under this experiment's rule, inner sweet starch-like pieces shift it to nachos.
watergate salad	yes	salad	Pudding-based dessert mixture with add-ins but no distinct structural starch shell, layer, or enclosed starch mass.

Table 1: Classifications for the mini experiment on foods whose names contain `salad`. *The narrower ontology question of whether candy-like inclusions in dishes such as Snickers salad should count as starch is discussed further in Section 6.2.

Theorem 1 (No Free Lunch Special)**.** *If the number of distinct foods assigned to tensor cells exceeds IJK, then at least one coordinate (i, j, k) must contain more than one food. Equivalently, if $|\mathbb{F}_{\text{real}}| > IJK$ (or if any subset of $\mathbb{F}_\infty$ of size $> IJK$ is mapped into the index space), then collisions are unavoidable.*

Proof (sketch). There are exactly IJK available cells in the tensor. Assigning more than IJK distinct foods to those cells forces at least one cell to receive two or more foods by the pigeonhole principle [2]. □

Example. A concrete collision can occur at $(i, j, k) = (\text{nachos}, \text{red_meat}, \text{pasta_noodle})$, where both pasta salad with ground beef and spaghetti with meatballs plausibly occupy the same cell. Such a case suggests that one might introduce a fourth dimension ℓ for temperature or a fifth dimension m for protein morphology if finer distinctions are desired.

6.3 Boundary Cases

The horseshoe theory of salad is an entertaining consequence of the current Cube Rule, but it also exposes a crack in the epistemological crust of the ontology. Mashed potatoes are treated as salad because they are a starch-based mono-food; by the same reasoning, a bowl of croutons appears to revert from nachos back to salad once the lettuce disappears. Yet the toast sandwich is canonically classified as sandwich rather than salad even though it, too, is effectively a mono-starch food, because its starch occupies top and bottom faces rather than none. Plain rice remains even less settled: it is the one food for which the Rule explicitly declines to decide.

Taken together, these cases suggest that morphology alone does not cleanly resolve mono-starch foods. Under the current ontology, one mono-starch becomes salad, another remains sandwich, and a third is left unclassified. This is

no longer just a funny edge case but a genuine modeling problem: the present nine categories do not distinguish between dishes whose ambiguity comes from mixed ingredients and dishes whose ambiguity comes from being composed of starch alone. In that sense the analogy to political horseshoe theory cuts twice: just as critics argue that horseshoe theory is reductive because it compresses emergent factions into an overly simple geometry, the Cube Rule also compresses mono-starches into categories built for more familiar morphologies rather than making room for them on their own terms [1]. The issue is not merely quirky; it remains slightly underdone at the center.

That tension motivates a tenth cube category reserved for *mono-starches* or *single-starch* foods. Such a category would be distinct from both salad (no structural starch) and nachos (outer container plus inner starch pieces), and would absorb plain rice, plain croutons, mashed potatoes, and related starch-based mono-foods without forcing them into categories whose logic was designed for other structures. Under that extension, the horseshoe theory would remain useful as a description of how the current Cube Rule behaves, while the new category would serve as a cleaner repair to the ontology itself.

Our current dish-level classification does not add that tenth category, so we preserve the Rule's open interpretation for mono-starch boundary cases in the reported tensor. Even that repair would leave at least one potato-shaped problem on the table: where should potato-filled pierogis go? Structurally they are calzone-like, since starch encloses the filling from the outside; by filling logic they also look nacho-like, since the interior is itself starch; and under a mono-starch extension they might be claimed as starch through and through. We leave that exercise to the reader. Relatedly, jello-style dishes suggest that gelatin may warrant a separate protein index j beyond the current formalism, likely yielding a very sparse but interesting subset of $\mathcal{S}$ in the morphological state space (with soup dumplings as one motivating example, since they classify as calzone under Cube Rule morphology and include gelatinized broth).

6.4 Topological Interpretation

Topological language is useful here: regions where $T_{i,j,k} = 1$ is dense in index space behave like low-curvature neighborhoods, while isolated $(i, j, k) \in \mathcal{S}$ resemble high-curvature outliers in a discrete morphological state space; for example, several bread-based categories may cluster while a more unusual dish such as jello salad sits farther from those neighborhoods. If local adjacency relations are imposed on occupied cells, $\mathcal{S}$ can also be viewed as inducing a sparse simplicial complex over the index lattice, with missing bridges and cavities corresponding to structurally absent combinations. Under that view, persistent homology and broader topological data analysis provide language for discussing connected components, holes, and other mesoscale features of the food-occupancy pattern as the ontology grows, such as whether sushi-like foods form one connected region while dessert salads occupy scattered outlier cells. Operationally, the ontology should avoid both excessive coarseness and excessive fragmentation: I, J, and K should be large enough to separate distinct dishes, but restrained enough to keep $|\mathcal{S}|$ empirically supportable.

7 Speculative Food Fiction

One practical use of the framework is imputation: once a neighborhood of the tensor is well populated, nearby empty cells become prompts for novel foods that are morphologically coherent even if they are not yet culturally established. The most productive region for this kind of food fiction is arguably the calzone morphology. If one asks, objectively, which morphology is most beloved, calzone has a strong claim: it contains dumplings, ravioli, empanadas, and other durable classics, while also absorbing newer handheld starch technologies such as Pop-Tarts and Uncrustables [3]. That pattern is not accidental. Fully enclosed starches travel well, protect fillings, and are easy to eat without utensils, so the calzone region of the tensor is both historically crowded and structurally hospitable to innovation.

Table 2: Accepted calzone candidates (`cube_type=calzone`) proposed by the LLM for cells treated as empty in the working catalog. Some may be rare or previously uncatalogued rather than genuinely novel.

Food	Cell (j, k)	Brief description
Beef Zongzi	`red_meat` `rice`	Glutinous rice dumpling with braised beef fully enclosed by rice.
Chicken Onigiri Ball	`poultry` `rice`	Compressed rice ball with a chicken center sealed on every surface.
Egg-Stuffed Potato Knish	`egg` `potato`	Potato-dough shell fully enclosing a seasoned egg filling.
Egg Piroshki	`egg` `wheat`	Sealed wheat-dough pocket stuffed with chopped egg and scallions.

This makes calzone a natural target for speculative completion. Given a desired protein-starch pairing, the framework can ask the LLM to impute a missing but morphologically valid enclosed dish, then accept only outputs whose descriptions preserve the six-face enclosure required by the Cube Rule. Table 2 shows a small sample of accepted calzone candidates proposed for cells treated as empty in the working catalog.

These examples are intentionally modest rather than fantastical. That modesty exposes an epistemic limitation in the setup: because $\mathbb{F}_{\text{real}}$ is a theoretical historical set that we only approximate with a finite catalog, an output proposed for a cell in $\mathcal{E}$ need not truly lie in $\mathbb{F}_\infty \setminus \mathbb{F}_{\text{real}}$. It may instead be a real but rare, poorly documented, regionally specific, or simply uncatalogued food that was missing from the menu of our current dataset. In other words, the pipeline is better understood as generating candidates at the frontier between the known support of the catalog and the larger, only partially observed set $\mathbb{F}_{\text{real}}$.

This raises a natural limiting question: as culinary innovation accumulates over long time scales, does $\mathbb{F}_{\text{real}}$ asymptotically approach $\mathbb{F}_\infty$, making genuinely empty cells increasingly scarce? We do not answer that here. But the framework remains useful under either view. If the gap $\mathbb{F}_\infty \setminus \mathbb{F}_{\text{real}}$ stays large, then the method is a tool for genuine novelty imputation; if the gap is small, then it is still a productive way to surface interesting, low-attestation, or culturally peripheral combinations that inhabit underexplored regions of the tensor. In that weaker but still valuable sense, speculative food fiction is less about inventing nonsense than about finding foods that the morphology suggests ought to exist, whether or not history has already gotten there first. The point is less to spice things up for its own sake than to probe which empty cells are structurally plausible.

8 Conclusion

We have shown that the Cube Rule of Food can be expressed as a geometric classification problem: dishes map to coordinates (i, j, k) in a third-order tensor T, with cube morphology i, protein taxonomy j, and starch substrate k as the axes of a discrete culinary state space. Canonical Cube Rule examples occupy well-defined cells; empty cells correspond to conceivable-but-unobserved dishes; and collisions—multiple foods at the same (i, j, k)—become topological facts rather than matters of taste. By constraining an LLM to propose candidates under strict axis matching, lexical guardrails, and a review gate, we have extended the support set $\mathcal{S}$ beyond the canonical seed while keeping the tensor falsifiable and the Rule's edge cases explicit: salad occupies the `starch_type=none` slice; the horseshoe theory of salad captures how the current Rule treats the transition from single-component salad (all lettuce or all croutons) to nachos (lettuce and croutons together); and mono-starch foods reveal a remaining ambiguity in the ontology rather than a fully resolved corner case. Methodologically, this combines prompt taxonomies, lightweight retrieval-augmented generation, guided decoding, and constrained tensor completion within a single food-classification pipeline, for example when canonical dishes such as pizza slice or nigiri are supplied as exemplars before the model proposes new dishes for vacant cells.

This formalism can also be extended in principled ways. The index sets I, J, and K need not

remain fixed. The cube axis i may gain a dedicated tenth *single-starch* category, resolving the present mismatch in which mashed potatoes fall toward salad, toast sandwich remains sandwich, and plain rice is left to the reader's interpretation. The protein axis j may admit gelatin (and thus jello, soup dumplings, and other aspic-bound dishes) as a distinct type, yielding a sparse but coherent region of $\mathcal{S}$. The starch axis k may be extended to include sugar-dominant substrates—candy, nougat, snickers bars—so that snickers salad and its ilk cease to be boundary cases and take their place at a named coordinate. Such an ontological change, however, should be made with care and with attention to cultural naming practices: doing so would create tension between the Midwestern term "snickers salad" and the strict starch-morphology Cube Rule definition of "salad," potentially opening taxonomic rifts between competing ontologies. Each such expansion refines the representation of conceivable foods within the same indexing framework.

More broadly, this framework treats morphology, protein, and starch as formal indices rather than informal descriptors. Future work can therefore extend the same tensor formalism while keeping the convention that salad occupies the `starch_type=none` slice and treating monostarch foods as an explicit ontological question rather than an incidental oddity.

Acknowledgements

This work would of course not have been possible without Phosphatide's original food-identification post and André Arko's later Cube Rule expansion and popularization. The author thanks Connor Wilson for inspiring this idea, Michael Davinroy for introducing the concepts of snickers salad and jello salad, and Andrew McNutt for first sharing the spirit of this venue.

References

[1] Horseshoe theory. `https://en.wikipedia.org/wiki/Horseshoe_theory`.

[2] Pigeonhole principle. `https://en.wikipedia.org/wiki/Pigeonhole_principle`.

[3] Andre Arko. The cube rule of food identification. `https://cuberule.com/`, 2017. Website.

[4] Phosphatide. Grand unified theory of food identification. `https://x.com/Phosphatide/status/974067376894328833`, 2018. Twitter/X post.

[5] theemmitt and Reddit contributors. Every country has a dumpling. `https://www.reddit.com/r/AskEurope/comments/jse5m2/every_country_has_a_dumpling/`, 2020. Reddit discussion in r/AskEurope.

[6] Wikipedia contributors. Toast sandwich. `https://en.wikipedia.org/wiki/Toast_sandwich`, 2026.

[7] Manuela Zangara. Smoked salmon millefeuille. `https://www.manusmenu.com/smoked-salmon-mille-feuille`, 2014. Recipe page.